

Chapter 4

Memory Design: System on Chip and Board based systems

4.1 Introduction

Memory design is the key to system design. The memory system is often the most costly (in terms of area or dies) part of the system and it largely determines the performance. Regardless of the processors and the interconnect, the application cannot be executed any faster than the memory system, which provides the instructions and the operands.

Memory design involves a number of considerations. The primary consideration is the application requirements: the operating system, the size and the variability of the application processes. This largely determines the size of memory and how the memory will be addressed: real or virtual. Figure 4.1 is an outline for memory design, while Table 4.1 compares the area for different memory technologies.

We start by looking at issues in SOC external and internal memories. We then examine scratchpad and cache memory to understand how they operate, and how they are designed. After that we consider the main memory problem, first the on-die memory and then the conventional DRAM design. As part of the design of large memory systems, we look at multiple memory modules, interleaving and memory system performance.

Table 4.2 shows the types of memory that can be integrated into an SOC design.

Memory technology	Rbe	kB per unit A
DRAM	0.05 to 0.1	1,800 – 3,600
SRAM	0.6	300
ROM/PROM	0.2 to 0.8+	225 – 900
eDRAM	0.15	1,200
flash: NAND	0.02	10,000

Table 4.1 Area comparison for different memory technology.




Figure 4.1 An outline for memory design.


Example. Required functionality can play a big role in achieving performance. Consider the differences between the two paths of Figure 4.1: the maximum functionality path and the restricted functionality path. The difference seems slight; whether the memory is off-die or on-die. The resulting performance difference can be great because of the long off-die access time. If the memory (application data and program) can be contained in an on-die memory, access time will be 3 to 10 cycles.

Off-die access times are an order of magnitude greater (30 to 100 cycles). To achieve the same performance, the off-die memory design must have an order of magnitude more cache, often split into multiple levels to meet access time requirements. Indeed, a cache bit can be 50 times larger than an on-die eDRAM bit (see Chapter 2 and Section 4.13). So the true cost of the larger cache required for off-die memory support may be 10 by 50 or 500 DRAM bits. If a memory system uses 10K rbe for cache to support an on-die memory, the die would require 100K rbe to support off-die memory. That 90K rbe difference could possibly accommodate 450K eDRAM bits.

4.2 Overview

4.2.1 SOC external memory: flash

Flash technology is a rapidly developing technology with improvements announced regularly. Flash is not really a memory replacement but probably better viewed as a disk replacement. However, in some circumstances and configurations it can serve the dual purpose of memory and non-volatile backup storage.

**Figure 4.2** The SOC memory model.

Format	(approx.) Size in mm	Weight in grams	Speed (read/write)	Typical applications
CF (compact flash)	36×43×3.3	11.4	22/18 MBps	Digital camera
SD (secure digital)	32×24× 2.1	2.0	22/18 MBps	Digital/video camera
Mini SD	20×20×1.2	1.0	22/18	Cell phone, GPS
Micro SD	15×11×0.7	0.5	22/15	Mini cell phone

Table 4.2 Some flash memory (NAND) package formats (2–128 GB size).

Technology	NOR	NAND
Bit density (kB per A)	1,000	10,000
Typical capacity	64 MB	1–4 GB (dies can be stacked to 32 GB)
Access time	20–50 ns	10 μ s
Transfer rate	150 MBps	300 MBps
Write time	300 μ s	200 μ s
Addressability	Word or Block	Block
Application	Program storage and limited data store	Disk replacement

Table 4.3 Comparison of flash memories.

Flash memory consists of an array of floating gate transistors. These transistors are similar to MOS transistors but with a two gate structure: a control gate and an insulated floating gate. Charge stored on the floating gate is trapped there, providing a non volatile storage. While the data can be rewritten, the current technology has a limited number of reliable rewrite cycles; usually less than a million. Since degradation with use can be a problem, error detection and correction are frequently implemented. While the density is excellent for semiconductor devices, the write cycle limitation generally restricts the usage to storing infrequently modified data, such as programs and large files.

There are two types of flash implementations: NOR and NAND. The NOR implementation is more flexible but the NAND provides significantly better bit density. Hybrid NOR/NAND implementations are also possible with the NOR array acting as a buffer to the larger NAND array. Table 4.3 provides a comparison of these implementations.

Flash memory cards come in various package formats; larger sizes are usually older (see Table 4.2). Small flash dies can be “stacked” with a SOC chip to present a single system/memory package. Flash die can also be stacked to create large (16 GB) single memory packages.

In current technology Flash usually is found in off die implementations. However there are a number of Flash variants that are specifically designed to be compatible with ordinary SOC technology. SONOS [19] is a non-volatile example, and Z-RAM [21] is a DRAM replacement example. Neither seems to suffer from rewrite cycle limitations. Z-RAM seems otherwise compatible with DRAM speeds while offering improved density. SONOS offers density but with slower access time than eDRAM.

4.2.2 SOC internal memory: placement

The most important and obvious factor in memory system design is the placement of the main memory: on-die (the same die as the processor) or off-die (on its own die or on a module with multiple die). As pointed out in Chapter 1, this factor distinguishes conventional workstation processors and application oriented board designs from SOC designs.

Item	Workstation Type	SOC single die	SOC Board based
Processor	Fastest available	Smaller, perhaps 4–6 times slower	as with SOC
Cache	2–3 levels very large (4–32MB)	simple, single level (256KB)	single level, multi-element
Memory Bus	Complex, slow pin limited	internal, wide, high bandwidth	Mix
Bus control	Complex timing and control	Simple internal	Mix
Memory	Very large (4+GB), limited bandwidth	Limited size (256MB) relatively fast	Specialized on board
Memory access time	20–30 ns	3–5 ns	Mix

Table 4.4 Comparing system design environments.

The design of the memory system is limited by two basic parameters that determine memory systems performance. The first is the access time. This is the time for a processor request to be transmitted to the memory system, access a datum, and return it back to the processor. Access time is largely a function of the physical parameters of the memory system—the physical distance between the processor and the memory system, or the bus delay, the chip delay, etc. The second parameter is *memory bandwidth*, the ability of the memory to respond to requests per unit time. Bandwidth is primarily determined by the way the physical memory system is organized—the number of independent memory arrays and the use of special sequential accessing modes.

The cache system must compensate for limits on memory access time and bandwidth.

The workstation processor, targeting high performance requires a very efficient memory system; a task made difficult by memory placement off-die. Table 4.4 compares the memory system design environments.

The workstation and board based memory design is clearly a greater challenge for designers. Special attention must be paid to the cache, which must make up for the memory placement difficulties.

4.2.3 The size of memory

As it will become obvious in this chapter, designing for large off die memory is the key problem in system board designs. So why not limit memory to sizes that could be incorporated on a die? In a virtual memory system we can still access large address spaces for applications. For workstations the application environment (represented by the operating system) has grown considerably (see Figure 4.3). As the environment continues to grow, so too does the working set or active pages of storage. This requires more real (physical) memory to hold a sufficient number of pages to avoid excessive page swapping, which can destroy performance. Board based systems face a slightly




Figure 4.3 Required disk space for several generations of Microsoft's Windows operating system.

different problem. Here the media based data sets are naturally very large and require large bandwidth from memory and substantial processing ability from the media processor. Board based systems have an advantage, however, as the access time is rarely a problem so long as the bandwidth requirements are met. How much memory can we put on a die? Well, that depends on the technology (feature size) and the required performance. Table 4.1 shows the area occupied for various technologies. The eDRAM size assumes a relatively large memory array (see later in this chapter). So, for example, in a 90nm technology we might expect to have about 12.3 kA per cm^2 or about 2 MB of eDRAM. Advancing circuit design and technology could significantly improve that but it does seem that about 64MB would be a limit, unless a compatible flash technology becomes available.

4.3 Scratchpads and Cache Memory

Smaller memories are almost always faster than larger memory, so it is useful to keep frequently used (or anticipated) instructions and data in a small, easily accessible (one cycle access) memory. If this memory is managed directly by the programmer, it is called a scratchpad memory; if it is managed by the hardware, it is called a cache.

Since management is a cumbersome process, most general-purpose computers use only cache memory. SOC, however, offers the potential of having the scratchpad alternative. Assuming that the application is well-known, the programmer can explicitly control data transfers in anticipation of use. Eliminating the cache control hardware offers additional area for larger scratchpad size; again improving performance.

SOC implements scratchpads usually for data and not for instructions, as simple caches

work well for instructions. Furthermore it is not worth the programming effort to directly manage instruction transfers.

The rest of this section treats the theory and experience of cache memory. Because there has been so much written about cache it is easy to forget the simpler and older scratchpad approach but with SOC sometimes the simple approach is best.

Caches work on the basis of the locality of program behavior [14]. There are three principles involved:

1. **Spatial Locality** — Given an access to a particular location in memory, there is a high probability that other accesses will be made to either that or neighboring locations within the lifetime of the program.
2. **Temporal Locality** — Given a sequence of references to n locations, there will be references into the same locations with high probability.
3. **Sequentiality** — Given that a reference has been made to location s it is likely that within the next few references there will be a reference to location of $s + 1$. This is a special case of spacial locality.

The cache designer must deal with the processor's accessing requirements on the one hand, and the memory system's requirements on the other. Effective cache designs balance these within cost constraints.

4.4 Basic Notions

Processor references contained in the cache are called cache hits. References not found in the cache are called cache misses. On a cache miss, the cache fetches the missing data from memory and places it in the cache. Usually the cache fetches an associated region of memory called the line. The line consists of one or more physical words accessed from a higher level cache or main memory. The physical word is the basic unit of access to the memory.

The processor–cache interface has a number of parameters. Those that directly affect processor performance (Figure 4.4) include:

1. Physical word — unit of transfer between processor and cache.
Typical physical word sizes:

2–4 bytes	— minimum, used in small core type processors.
8 bytes and larger	— multiple instruction issue processors (superscalar).
2. Block size (sometimes called *line*) — usually the basic unit of transfer between cache and memory. It consists of n physical words transferred from main memory via the bus.
3. Access time for a cache hit — this is a property of the cache size and organization.
4. Access time for a cache miss — property of the memory and bus.




Figure 4.4 Parameters affecting processor performance.

5. Time to compute a real address given a virtual address (not-in-TLB time) — property of the address translation facility.
6. Number of processor requests per cycle.

Cache performance is measured by the miss rate or the probability that a reference made to the cache is not found. The miss rate times the miss time is the delay penalty due to the cache miss. In simple processors, the processor stalls on a cache miss.

Is cache part of the processor?

For many IP designs, the first level cache is integrated into the processor design, so what and why do we need to know cache details? The most obvious answer is that an SOC consists of multiple processors which must share memory, usually through a second level cache. Moreover, the details of the first level cache may be essential in achieving memory consistency and proper program operation. So for our purpose the cache is a separate, important piece of the SOC. We design the SOC memory hierarchy, not an isolated cache.

4.5 Cache Organization

A cache uses either a fetch on demand or a prefetch strategy. The former organization is widely used with simple processors. A demand fetch cache brings a new memory locality into the cache only when a miss occurs. The prefetch cache attempts to anticipate the locality about to be requested and *prefetches* it. It is commonly used in I-caches.

There are three basic types of cache organization: (fully) associative mapped (Figure 4.5), direct-mapped (Figure 4.6), and set associative mapped (Figure 4.7, which is really a combination of the other two). In a fully associative cache, when a request is




Figure 4.5 Fully associative mapping.

made, the address is compared with the addresses of all entries in the directory. If the requested address is found (a directory hit), the corresponding location in the cache is fetched; otherwise, a *miss* occurs.

In a direct-mapped cache, the lower order line address bits access the directory (index bits in Figure 4.8). Since multiple line addresses map into the same location in the cache directory, the upper line address bits (tag bits) must be compared with the directory address to validate a hit. If a comparison is not valid, the result is a miss. The advantage of the direct mapped cache is that a reference to the cache array itself can be made *simultaneously with the access to the directory*.

The address given to the cache by the processor is divided into several pieces, each of which has a different role in accessing data. In an address partitioned as in Figure 4.8, the most significant bits that are used for comparison (with the upper portion of a line address contained in the directory) are called the *tag*.

The next field is called the *index* and it contains the bits used to address a line in the cache directory. The *tag* plus the *index* is the line address in memory.

The next field is the *offset* and it is the address of a physical word within a line.





Figure 4.6 Direct mapping.

Finally, the least significant address field specifies a *byte in a word*. These bits are not usually used by the cache, since the cache references a word. (An exception arises in the case of a *write* that modifies only a part of a word.)

The set associative cache is similar to the direct-mapped cache. Bits from the line address are used to address a cache directory. However, now there are multiple choices: two, four, or more complete line addresses may be present in the directory. Each address corresponds to a location in a sub-cache. The collection of these sub-caches forms the total cache array. These sub-arrays can be accessed simultaneously, together with the cache directory. If any of the entries in the cache directory match the reference address there is a hit, that sub-cache array is outgoing back to the processor. While selection in the outgoing process increases the cache access time, the set associative cache access time is usually better than that of the associative mapped cache. But the direct-mapped cache provides the fastest processor access to cache data for any given cache size.

4.6 Cache Data

Cache size largely determines cache performance (miss rate). The larger the cache, the lower the miss rate. Almost all cache miss rate data is empirical, and as such has certain limitations. Cache data is strongly program dependent. Also data is frequently based upon older machines, where the memory and program size was fixed and small. Such data shows low miss rate for relatively small size caches. Thus, there is a tendency

**Figure 4.7** Set associative (multiple direct-mapped caches).**Figure 4.8** Address partitioned by cache usage.




Figure 4.9 A design target miss rate per reference to memory (fully associative, demand fetch, fetch (allocate) on write, copyback with LRU replacement) [17], [18].

for the measured miss rate of a particular cache size to increase over the time. This is simply the result of measurements made on programs of increasing size. Some time ago Smith [18] developed a series of design target miss rates (DTMR) that represent an estimate of what a designer could expect from an integrated (instruction and data) cache. These data are presented in Figure 4.9 and gives an idea of typical miss rates as a function of cache and line sizes.

For cache sizes larger than 1MB, a general rule is that doubling the size halves the miss rate. The general rule is less valid in transaction based programs.

4.7 Write Policies

How is memory updated on a write? One could write to both cache and memory (write-through), or write only to the cache (copyback - sometimes called write back), updating memory when the line is replaced. These two strategies are the basic cache write policies (Figure 4.10).

The write-through cache (Figure 4.10a) stores into both cache and main memory on each CPU store.


Advantage: this retains a consistent (up-to-date) image of program activity in memory.

Disadvantage: memory bandwidth may be high—dominated by write traffic.

In the copyback cache (Figure 4.10b), the new data is written to memory when the line is replaced. This requires keeping track of modified (or "dirty") lines, but results in reduced memory traffic for writes.

1. Dirty bit is set if a write occurs anywhere in line.

a) write through



b) write back (copy back)




Figure 4.10 Write policies. (a) Write-through cache (no allocate on write); (b) Copy-back cache (allocate on write).

2. From various traces [17], the probability that a line to be replaced is dirty is 47% on average (ranging from 22% to 80%).
3. Rule of thumb: half of the data lines replaced are dirty. So, for a data cache assume 50% are dirty lines and for an integrated cache assume 30% are dirty lines.

Most larger caches use copyback; write-through is usually restricted to either small caches or special purpose caches that provide an up to date image of memory.

4.8 Strategies for Line Replacement at Miss Time

What happens on a cache miss? If the reference address is not found in the directory, a *cache miss* occurs. Two actions must promptly be taken: (1) the missed line must be fetched from main memory; (2) one of the current cache lines must be designated for replacement by the currently accessed line (the missed line).

4.8.1 Fetching a Line

In a write-through cache fetching a line involves accessing the missed line and the replaced line is discarded (written over).

For a copyback policy, we first determine whether the line to be replaced is *dirty* (has been written to) or not. If the line is clean, the situation is the same as with the write-

through cache. However, if the line is dirty, we must write the replaced line back to memory.

In accessing a line, the fastest approach is the *nonblocking* cache or *prefetching* cache. This approach is applicable in both write-through and copyback caches. Here the cache has additional control hardware to allow the cache miss to be handled (or bypassed) while the processor continues to execute. This strategy only works when the miss is accessing cache data that is not currently required by the processor. Nonblocking caches perform best with compilers that provide prefetching of lines in anticipation of processor use. The effectiveness of nonblocking caches depends on:

1. The number of misses that can be bypassed while the processor continues to execute.
2. The effectiveness of the prefetch and the adequateness of the buffers to hold the prefetch information. The longer the prefetch is made before expected use, the less the miss delay; but this also means that the buffers or registers are occupied and hence not available for (possible) current use.

4.8.2 Line Replacement

The replacement policy selects a line for replacement when the cache is full. There are three replacement policies that are in wide use:

1. Least Recently Used (LRU) — The line that was least recently accessed (by a read or write) is replaced.
2. First In–First Out (FIFO) — The line that had been in the cache the longest is replaced.
3. Random Replacement (RAND) — Replacement is determined randomly.

Since the LRU policy corresponds to the concept of temporal locality it is generally the preferred policy. It is also the most complex to implement. Each line has a counter which is updated on a read (or write). Since these counters could be large, it is common to create an approximation to the true LRU with smaller counters.

While LRU performs better than either FIFO or RAND, use of the simpler RAND or FIFO only amplifies the LRU miss rate (DTMR) by about 1.10 (i.e., 10%) [17].

4.8.3 Cache Environment: Effects of System, Transactions and Multiprogramming

Most available cache data is based upon trace studies of user applications. Actual applications are run in the context of the system. The operating system tends to slightly increase (20% or so) the miss rate experienced by a user program [10].

Multiprogramming environments create special demands on a cache. In such environments the cache miss rates may not be affected by increasing cache size. There are two environments:




Figure 4.11 Warm cache: cache miss rates for a multiprogrammed environment switching processes after Q instructions.

1. A multiprogrammed environment. The system together with several programs is resident in memory. Control is passed from program to program after a number of instructions, Q, have been executed, and eventually returns to the first program. This kind of environment results in what is called a *warm cache*. When a process returns for continuing execution, it finds that some, but not all, of its most recently used lines in the cache, increasing the expected miss rate (Figure 4.11 illustrates the effect).
2. Transaction processing. While the system is resident in memory together with a number of support programs, short applications (transactions) run through to completion. Each application consists of Q instructions. This kind of environment is sometimes called a *cold cache*. Figure 4.12 illustrates the situation.

Both of the preceding environments are characterized by passing control from one program to another before completely loading the working set of the program. This can significantly increase the miss rate.

4.9 Other Types of Cache

So far we have considered only the simple integrated cache (also called a “unified” cache) which contains both data and instructions. In the next few sections we consider various other types of cache. The list we present (Table 4.5) is hardly exhaustive, but it illustrates some of the variety of cache designs possible for special or even commonplace applications.

Most currently available microprocessors use split I/D caches, described in the next section.




Figure 4.12 Cold cache: cache miss rates for a transaction environment switching processes after Q instructions.

Type	Where Usually Used
Integrated (or “unified”)	The basic cache which accommodates all references (I and D). This is commonly used as the second and higher level cache.
Split cache I and D	Provides additional cache access bandwidth at some loss of MR. Commonly used as a first level processor cache.
Sectored cache	Improves area effectiveness (MR for given area) for on-chip cache.
Multi-level cache	First level has fast access, second level is usually much larger than the first to reduce time delay in a first-level miss.
Write assembly cache	Specialized; reduces write traffic; usually used with a WT on-chip first-level cache.

Table 4.5 Common types of cache.

4.10 Split I- and D-Caches and the Effect of Code Density

Multiple caches can be incorporated into a single processor design, each cache serving a designated process or use. Over the years, special caches for systems code and user code or even special I/O caches have been considered. The most popular configuration of partitioned caches is the use of separate caches for instructions and data.

Separate instruction and data caches provides significantly increased cache bandwidth, doubling the access capability of the cache ensemble. I and D caches come at some expense, however; a unified cache with the same size as the sum of a data and instruction cache has a lower miss rate. In the unified cache, the ratio of instruction to data working set elements changes during the execution of the program and is adapted to by the replacement strategy.

Split caches have implementation advantages. Since the caches need not be split equally, a 75%–25% or other split may prove more effective. Also, the I-cache is simpler as it is not required to handle stores.

4.11 Multi-Level Caches

4.11.1 Limits on cache array size

The cache consists of an SRAM array of storage cells. As the array increases in size so does the length of the wires required to access its most remote cell. This translates into the cache access delay which is a function of both the cache size, organization and the process technology (feature size, f). McFarland [1] has modeled the delay and found that an approximation can be represented as:

$$\text{Access time (ns)} = (0.35 + 3.8f + (0.006 + 0.025f)C) \times (1 + 0.3(1 - 1/A))$$

where f is the feature size in microns, C is the cache array capacity in KB, and A is the degree of associativity (where direct map = 1).

The effect of this equation (for $A = 1$) can be seen in Figure 4.13. If we limit the level 1 access time to under 1ns, we are probably limited to a cache array of about 32KB. While it's possible to interleave multiple arrays, the interleaving itself has an overhead. So usually L1 caches are less than 64KB, L2 caches are usually less than 512KB (probably interleaved using smaller arrays) and L3 caches use multiple arrays of 256KB to create large caches, often limited by die size.

4.11.2 Evaluating multi-level caches

In the case of a multi level cache we can evaluate performance of both levels using L1 cache data. A two-level cache system is termed *inclusive* if all the contents of the lower-level cache (L1) are also contained in the higher-level cache (L2).

Second-level cache analysis is achieved using the principle of inclusion—i.e., a large second-level cache includes the same lines as in the first-level cache. Thus for the purpose of evaluating performance, we can assume that the first-level cache does not exist.





Figure 4.13 Cache access time (for a single array) as a function of cache array size.

The total number of misses that occur in a second-level cache can be determined by assuming that the processor made all of its requests to the second-level cache without the intermediary first-level cache.

There are design considerations in accommodating a second-level cache to existing first-level cache. The line size of the second-level cache should be the same as or larger than the first-level cache. Otherwise, if the line size in the second-level cache were smaller, loading the line in the first-level cache would simply cause two misses in the second-level cache. Further, the second-level cache should be significantly larger than the first-level otherwise it would have no benefit.

In a two-level system as shown in Figure 4.14 with first-level cache, L1, and second-level cache, L2, we define the following miss rates [15]:


1. A local miss rate is simply the number of misses experienced by the cache divided by the number of references to it. This is the usual understanding of *miss rate*.
2. The global miss rate is the number of L2 misses divided by the number of references made by the processor. This is our primary measure of the L2 cache.
3. The solo miss rate is the miss rate the L2 cache would have if it were the only cache in the system. This is the miss rate defined by the principle of inclusion. If L2 contains *all* of L1 then we can find the number of L2 misses and the processor reference rate, ignoring the presence of the L1 cache. The principle of inclusion specifies that the global miss rate is the same as the solo miss rate, allowing us to use the solo miss rate to evaluate a design.

**Figure 4.14** A two-level cache.

The preceding data (read misses only) illustrate some salient points in multi-level cache analysis and design:

1. So long as the L1 cache is the same as or larger than the L2, analysis by the principle of inclusion provides a good estimate of the behavior of the L2 cache.
2. When the L2 cache is significantly larger than the L1 it can be considered independent of the L1 parameters. Its miss rate corresponds to a solo miss rate.

EXAMPLE 4.1



Miss penalties:

Miss in L1, hit in L2: 2 cycles
Miss in L1, miss in L2: 15 cycles

Suppose we have a two-level cache with miss rates of 4% (L1) and 1% (L2). Suppose the miss in L1 and hit in L2 penalty is 2 cycles, and the miss penalty in both caches is 15 cycles (13 cycles more than a hit in L2). If a processor makes 1 references per instruction, we can compute the excess CPI due to cache misses as follows:

SOC	L1 Cache	L2 Cache
NetSilicon NS9775 [9]	8kB I-cache, 4kB D-cache	-
NXP LH7A404 [7]	8kB I-cache, 8kB D-cache	-
Freescale e600 [3]	32kB I-cache, 32kB D-cache	1MB with ECC
Freescale PowerQUICC III [4]	32kB I-cache, 32kB D-cache	256kB with ECC
ARM1136J(F)-S [2]	64kB I-cache, 64kB D-cache	max 512kB

Table 4.6 SOC cache organization.

Excess CPI due to L1 misses

$$\begin{aligned} &= 1.0 \text{ refr/inst} \times 0.04 \text{ misses/refr} \times 2 \text{ cycles/miss} \\ &= 0.08 \text{ CPI} \end{aligned}$$

Excess CPI due to L2 misses

$$\begin{aligned} &= 1.0 \text{ refr/inst} \times 0.01 \text{ misses/refr} \times 13 \text{ cycles/miss} \\ &= 0.13 \text{ CPI} \end{aligned}$$

Note: the L2 miss penalty is 13 cycles, not 15 cycles, since the 1% L2 misses have already been “charged” 2 cycles in the excess L1 CPI.

$$\begin{aligned} \text{Total effect} &= \text{excess L1 CPI} + \text{excess L2 CPI} \\ &= 0.08 + 0.13 \\ &= 0.21 \text{ CPI} \end{aligned}$$




4.11.3 Logical Inclusion

True or logical inclusion, where *all* the contents of L1 reside also in L2, should not be confused with statistical inclusion, where *usually* L2 contains the L1 data. There are a number of requirements for logical inclusion. Clearly, the L1 cache must be write-through; the L2 cache needs not be. If L1 were copyback, then a write to a line in L1 would not go immediately to L2, so L1 and L2 would differ in contents.

When logical inclusion is required, it is probably necessary to actively force the contents to be the same and to use consistent cache policies.

Logical inclusion is a primary concern in shared memory multiprocessor systems that require a consistent memory image.

SOC	Organization	Number of entries
Freescale e600 [3]	Separate I-TLB, D-TLB	128 entry, 2-way set-associative, LRU
NXP LH7A404 [7]	Separate I-TLB, D-TLB	64 entry each
NetSilicon NS9775 (ARM926EJ-S) [9]	Mixed	32 entry 2-way set-associative

Table 4.7 SOC TLB organization.**Figure 4.15** TLB with two-way set associativity.

4.12 Virtual-to-Real Translation

The translation lookaside buffer (TLB) provides the real addresses used by the cache by translating the virtual addresses into real addresses.

Figure 4.15 shows a two-way set associative TLB. The page address (the upper bits of the virtual address) is composed of the bits that require translation. Selected virtual address bits address the TLB entries. These are selected (or *hashed*) from the bits of the virtual address. This avoids too many address collisions, as might occur when both address and data pages have the same, say, “000,” low-order page addresses. The size of the virtual address index is equal to $\log_2 t$, where t is the number of entries in the TLB divided by the degree of set associativity. When a TLB entry is accessed, a virtual and real translation pair from each entry is accessed. The virtual addresses are compared to the virtual address tag (the virtual address bits that were not used in the index). If a match is found the corresponding real address is multiplexed to the output




Figure 4.16 Not-in-TLB rate.

of the TLB.

With careful assignment of page addresses, the TLB access can occur at the same time as the cache access. When a translation is not found in the TLB, the process described in Chapter 1 must be repeated to create a correct virtual-to-real address pair in the TLB. This may require more than 10 cycles; TLB misses—called not-in-TLB—are costly to performance. TLB access in many ways resembles cache access. Fully associative organization of TLB is generally slow, but 4 way or higher set associative TLBs perform well and are generally preferred.

Typical TLB miss rates are shown in Figure 4.16. Fully associative (FA) data is similar to 4 way set associative.

For those SOC or board based systems that use virtual addressing there are additional considerations:

1. Small TLBs may create excess not-in-TLB faults, adding time to program execution.
2. If the cache uses real addresses, the TLB access must occur before the cache access, increasing the cache access time.

Excess not-in-TLB translations can generally be controlled through the use of a well-designed TLB. The size and organization of the TLB depends on performance targets.

Typically, separate instruction and data translation lookaside buffers (TLBs) are used. Both TLBs might use 128-entry, two-way set-associative, and use LRU replacement algorithm.




Figure 4.17 On-die SRAM and DRAM. The eDRAM must accommodate the process requirements of the logic, representing an overhead. SRAM is unaffected.

4.13 SOC (on die) Memory Systems

On-die memory design is a special case of the general memory system design problem, considered in the next section. The designer has much greater flexibility in the selection of the memory itself and the overall cache-memory organization. Since the application is known the general size of both the program and data store can be estimated. Frequently part of the program store is designed in a fixed ROM. The remainder of memory is realized with either static RAM (SRAM) or dynamic RAM (DRAM). While the SRAM is realized in the same process technology as the processor, DRAM is not. An SRAM bit consists of a 6 transistor cell, while the DRAM uses only one transistor plus a deep trench capacitor. The DRAM cell is designed for density, it uses few wiring layers. DRAM design target low refresh rates and hence low leakage currents. A DRAM cell uses a non minimum length transistor with higher threshold voltage, (V_T), to provide lower leakage current. This leads to lower gate overdrive and slower switching. For a standalone die the result is that the SRAM is 10–20 times faster and 10 or more times less dense than DRAM.

Embedded DRAM (eDRAM) [13, 11] has been introduced as a compromise for use as an on die memory. Since there are additional process steps in realizing an SOC with eDRAM, the macro to generate the eDRAM is fabrication specific and is regarded as a hard (or at least firm) IP. The eDRAM has an overhead (Figure 4.17) resulting in less density than DRAM. Process complexity for the eDRAM can include generating 3 additional mask layers resulting in 20% additional cost than that for the DRAM.

An SOC, using the eDRAM approach, integrates high speed, high leakage logic transistors with lower speed, lower leakage memory transistors on the same die. The advantage for eDRAM lies in its density as shown in Figure 4.18. Therefore, one key factor for selecting eDRAM over SRAM is the size of the memory required.

Having paid the process costs for eDRAM, the timing parameters for eDRAM are much better than conventional DRAM. The cycle time (and access time) are much closer to SRAM, as shown in Figure 4.19. All types of on die memory enjoy the advantage of bandwidth as a whole memory column can be accessed each cycle.

A final consideration in memory selection is the projected error rate due to radiation (called the *soft error rate* or SER). Each DRAM cell stores significantly larger amounts




Figure 4.18 The relative density advantage eDRAM improves with memory size.




Figure 4.19 Cycle time for random memory accesses.




Figure 4.20 The soft error ratio of SRAM to eDRAM.

of charge than in the SRAM cell. The SRAM cells are faster and easier to flip, with correspondingly higher soft error rate (SER). Additionally, for an SRAM cell as technology scales, the critical amount of charge for determining an error decreases due to scaling of supply voltages and cell capacitances. The differences are shown in Figure 4.20. At even 130 nm feature size the SER for SRAM is about 1800 times higher than for eDRAM. Of course, more error prone SRAM implementation can compensate by more extensive use of error correcting codes, but this comes with its own cost.

Ultimately the selection of on die memory technology depends on fabrication process access and memory size required.

4.14 Board based (off-die) Memory Systems

In many processor design situations (probably all but the SOC case), the main memory system is the principal design challenge.

As processor ensembles can be quite complex, the memory system that serves these processors is correspondingly complex.

The memory module consists of all the memory chips needed to forward a cache line to the processor via the bus. The cache line is transmitted as a burst of bus word transfers. Each memory module has two important parameters: module access time and module cycle time. The module access time is simply the amount of time required to retrieve a word into the output memory buffer register of a particular memory module, given a valid address in its address register. Memory service (or cycle) time is the minimum time between requests directed at the same module. Various technologies present a significant range of relationships between the access time and the service time. The access time is the total time for the processor to access a line in memory. In a small, simple memory system this may be little more than chip access time plus




Figure 4.21 Accessing delay in a complex memory system. Access time includes chip accessing, module overhead, bus transit, etc.

some multiplexing and transit delays. The service time is approximately the same as the chip cycle time. In a large, multi-module memory system (Figure 4.21), the access time may be greatly increased, as it now includes the module access time plus transit time, bus accessing overhead, error detection and correction delay, etc.

After years of rather static evolution of DRAM memory chips, recent years have brought about significant new emphasis on the performance (rather than simply the size) of memory chips.

The first major improvement to DRAM technology is synchronous DRAM or SDRAM. This approach synchronizes the DRAM access and cycle time to the bus cycle. Additional enhancements accelerate the data transfer and improves the electrical characteristics of the bus and module. There are now multiple types of SDRAM. The basic DRAM types are:

1. DRAM: Asynchronous DRAM
2. SDRAM: Synchronous DRAM. The memory module timing is synchronized to the memory bus clock.
3. DDR SDRAM: Double data rate SDRAM. The memory module fetches a double sized transfer unit each bus cycle and transmits at twice the bus clock rate.

In the next section we present the basics of the (asynchronous) DRAM, following that the more advanced SDRAMs.

4.15 Simple DRAM and the memory array

The simplest asynchronous DRAM consists of a single memory array with 1 (and sometimes 4 or 16) output bits. Internal to the chip is a two-dimensional array of memory cells consisting of rows and columns. Thus, half of the memory address is used to specify a row address, one of $2^{n/2}$ row lines, and the other half of the address is similarly used to specify one of $2^{n/2}$ column lines (Figure 4.22). The cell itself that holds the data is quite simple, consisting merely of an MOS transistor holding a charge




Figure 4.22 A memory chip.

(a capacitance). As this discharges over time, it must continually be refreshed, once every several milliseconds.

With large-sized memories, the number of address lines dominates the pinout of the chip. In order to conserve these pins and provide a smaller package for better overall density, the row and column addresses are multiplexed onto the same lines (input pins) for entry onto the chip. Two additional lines are important here: RAS (row address strobe) and CAS (column address strobe). These gate first the row address, then the column address into the chip. The row and column addresses are then decoded to select one out of the $2^{n/2}$ possible lines. The intersection of the active row and column lines is the desired bit of information. The column line's signals are then amplified by a sense amplifier and transmitted to the output pin (data out, or D_{out}) during a read cycle. During a write cycle, the write-enable signal stores the data-in (D_{in}) signal to specify the contents of the selected bit address.

All of these actions happen in a sequence approximated in the timing diagram in Figure 4.23. At the beginning of a read from memory, the RAS line is activated. With the RAS active and the CAS inactive, the information on the address lines is interpreted as the row address and stored into the row address register. This activates the row decoder and the selected row line in the memory array. The CAS is then activated, which gates the column address lines into a column address register. Note that:

1. The two rise times on CAS represent the earliest and latest that this signal may rise with respect to the column address signals.
2. WE is write enable. This signal is inactive during read operations.

The column address decoder then selects a column line; at the intersection of the row and column line is the desired data bit. During a read cycle the write enable is inactive




Figure 4.23 Asynchronous DRAM chip timing.

(low) and the output line (D_{out}) is at a high-impedance state until it is activated either high or low depending on the contents of the selected memory cell.

The time from the beginning of RAS until the data output line is activated is a very important parameter in the memory module design. This is called the chip access time or $t_{\text{chip access}}$. The other important chip timing parameter is the cycle time of the memory chip ($t_{\text{chip cycle}}$). This is not the same as the access time, as the selected row and column lines must recover before the next address can be entered and the read process repeated.

The asynchronous DRAM module does not simply consist of memory chips (Figure 4.24). In a memory system with p bits per physical word, and 2^n words in a module, the n address bits enter the module and are usually directed at a Dynamic Memory Controller chip. This chip in conjunction with a Memory Timing Controller provides the following functions:

1. Multiplexing the n address bits into a row and a column address for use by the memory chips.
2. The creation of the correct RAS and CAS signal lines at the appropriate time.
3. Provide timely refresh of the memory system.

Since the Dynamic Memory Controller output drives all p bits, and hence p chips, of the physical word, the controller output may also require buffering. As the memory





Figure 4.24 An asynchronous DRAM memory module.

read operation is completed, the data-out signals are directed at bus drivers which then interface to the memory bus, which is the interface for all of the memory modules.

Two features found on DRAM chips affect the design of the memory system. These "burst" mode type features are called:

1. Nibble mode.
2. Page mode.

Both of these are techniques for improving the transfer rate of memory words. In nibble mode, a single address (row and column) is presented to the memory chip and the CAS line is toggled repeatedly. Internally, the chip interprets this CAS toggling as a $\text{mod } 2^w$ progression of low-order column addresses. Thus, sequential words can be accessed at a higher rate from the memory chip. For example, for $w = 2$, we could access four consecutive low-order bit addresses, e.g.:



and then return to the original bit address.

In page mode, a single row is selected and nonsequential column addresses may be entered at a high rate by repeatedly activating the CAS line (similar to nibble mode, Figure 4.23). Usually this is used to fill a cache line.

While terminology varies, nibble mode usually refers to the access of (up to) four consecutive words (a nibble) starting on a quad word address boundary.

SOC	Memory type	Memory size	Memory position
Intel PXA27x [5]	SRAM	256KB	on-die
Philips Nexperia PNX1700 [6]	DDR SDRAM	256MB	off-die
Intel IOP333 [8]	DDR333	2GB	off-die

Table 4.8 SOC memory designs.




Figure 4.25 Internal configuration of DDR SDRAM.

4.15.1 SDRAM and DDR SDRAM

The first major improvement to the DRAM technology is the synchronous DRAM or the SDRAM. This approach, as mentioned before, synchronizes the DRAM access and cycle to the bus cycle. This has a number of significant advantages. It eliminates the need for separate memory clocking chips to produce the RAS and CAS signals. The rising edge of the bus clock provides the synchronization. Also by extending the package to accommodate multiple output pins, versions that have 4, 8 and 16 pins allow more modularity in memory sizing.

With the focus on the bus and memory-bus interface, we further improve bus bandwidth by using differential data and address lines. Now when the clock line rises, the complement clock falls, but midway through the cycle the clock line falls and the complement clock rises. This affords the possibility to transmit synchronous data twice during each cycle: once on the rising edge of the clock signal, and once on the rising




Figure 4.26 A line fetch in DDR SDRAM.

edge of the complement clock. Using this we are able to double the data rate transmitted on the bus. The resulting memory chips are called double data rate SDRAMs or DDR SDRAMs (Figure 4.25). Also instead of selecting a row and a column for each memory reference, it is possible to select a row and leave it selected (active) while multiple column references are made to the same row (Figure 4.26).

In cases where spatial locality permits, the read and write times are improved by eliminating the row select delay. Of course when a new row is referenced then the row activation time must be added to the access time. Another improvement introduced in SDRAMs is the use of multiple DRAM arrays, usually either 4 or 8. Depending on the chip implementation these multiple arrays can be independently accessed or sequentially accessed, as programmed by the user. In the former case each array can have an independently activated row providing an interleaved access to multiple column addresses. If the arrays are sequentially accessed then the corresponding rows in each array are activated and longer bursts of consecutive data can be supported. This is particularly valuable for graphics applications.

The improved timing parameters of the modern memory chip results from careful attention to the electrical characteristic of the bus and the chip. In addition to the use of differential signaling (initially for data, now for all signals), the bus is designed to be a terminated strip transmission line. With the DDR3 (now commonly called GDDR–graphics DDR) the termination is on die (rather than simply at the end of the bus) and special calibration techniques are used to ensure accurate termination.

The DDR chips that support interleaved row accesses with independent arrays must carry out a $2n$ data fetch from the array to support the double data rate. So a chip with 4 data out ($n = 4$) lines must have arrays that fetch 8 bits. The DDR2 arrays typically fetch $4n$, so with $n = 4$, the array would fetch 16 bits. This enables higher data transmission rates as the array is accessed only once for every 4 bus half-cycles.

Some typical parameters are shown in Table 4.9 and Table 4.10. While representatives of all of these DRAMs are in production at the time of writing, the asynchronous DRAM and the SDRAM are legacy parts and are generally not used for new development. The DDR3 part was introduced for graphic application configurations. For

	DRAM	SDRAM	DDR1	DDR2	GDDR (DDR3)
Typical chip capacity		1 Gb	1Gb	1Gb	256Mb
Output data pins/chip	1	4	4	4	32
Array data bits fetched	1	4	8	16	32
Number of arrays	1	4	4	8	4
Number of chips/module	64+	16	16	16	4
Burst word transfers	1-4	1,2,4	2,4,8	4,8	4,8,16
Rows			16k	16k	
Columns			2048x8	512x16	512x32
32B lines/ row/ array			2048	1024	2048x4

Table 4.9 Configuration parameters for some typical DRAM chips used in a 64bit module

	DRAM	SDRAM	DDR1	DDR2	DDR3
Bus clock rate (MHz)	Asynchronous	100	133	266	600
ACTIVE to CAS, ns	30	30	20	15	12
Col. addr. to data out1 (read time), ns	40	30	20	15	12
Line access, ns (accessing a new row)	140	90	51	36	28
Line access, ns (within an active row)	120	60	31	21	16
Rows interleaving	x1	x4	x4	x8	x1

Table 4.10 Timing parameters for some typical DRAM modules (64b)




Figure 4.27 SDRAM channels and controller

most cases the parameters are typical and for common configurations. For example the asynchronous DRAM is available with 1,4, and 16 output pins. The DDR SDRAMs are available with 4, 8 and 16 output pins. Many other arrangements are possible.

Multiple (up to 4) DDR2 SDRAMs can be configured to share a common bus (Figure 4.27). In this case when a chip is “active” (i.e., it has an active row) the on die termination is unused. When there are no active rows on the die the termination is used. Typical server configurations then might have 4 modules sharing a bus (called a channel) and a memory controller managing up to two busses (see Figure 4.27). The limit of two is caused simply by the large number of tuned strip and micro strip transmission lines that must connect the controller to the busses. More advanced techniques place a channel buffer between the module and a very high speed channel. This advanced channel has a smaller width (e.g. 1 byte) but a much higher data rate (e.g. 8×). The net effect leaves the bandwidth per module the same but now the number of wires entering the controller has decreased enabling the controller to manage more channels (e.g. 8).

4.15.2 Memory Buffers

The processor can sustain only a limited number of outstanding memory references before it suspends processing and the generation of further memory references. This can happen either as a result of logical dependencies in the program or because of an insufficient buffering capability for outstanding requests. The significance of this is that the achievable memory bandwidth is decreased as a consequence of the pause in the processing, for the memory can service only as many requests as are made by the processor.

Examples of logical dependencies include branches and address interlocks. The program must suspend computation until an item has been retrieved from memory.

Associated with each outstanding memory request is certain information that specifies the nature of the request (e.g. a read or a write operation), the address of the memory

location, and sufficient information to route requested data back to the requestor. All this information must be buffered either in the processor or in the memory system until the memory reference is complete. When the buffer is full, further requests cannot be accepted requiring the processor be stalled.

In interleaved memory, the modules usually are not all equally congested. So it is useful to maximize the number of requests made by the processor, in the hope that the additional references will be to relatively idle modules and will lead to a net increase in the achieved bandwidth. If maximizing the bandwidth of memory is a primary objective, we need buffering of memory requests up to the point at which the logical dependencies in the program become the limiting factor.

4.16 Models of Simple Processor–Memory Interaction

In systems with multiple processors or with complex single processors, requests may congest the memory system. Either multiple requests may occur at the same time, providing bus or network congestion, or requests arising from different sources may request access to the memory system. Requests that cannot be immediately honored by the memory system result in memory systems contention. This contention degrades the bandwidth, and is possible to achieve from the memory system.

In the simplest possible arrangement, a single simple processor makes a request to a single memory module. The processor ceases activity (as with a blocking cache) and waits for service from the module. When the module responds, the processor resumes activity. Under such an arrangement the results are completely predictable. There can be no contention of the memory system, since only one request is made at a time to the memory module. Now suppose we arrange to have n simple processors access m independent modules. Contention develops when multiple processors access the same module. Contention results in a reduced average bandwidth available to each of the processors. Asymptotically, a processor with a non-blocking cache making n requests to the memory system during a memory cycle resembles the n processor m module memory system, at least from a modeling point of view. But in modern systems, processors are usually buffered from the memory system. Whether or not a processor is slowed down by memory or bus contention during cache access depends on the cache design and the service rate of processors that share the same memory system.

Given a collection of m modules each with service time T_c , access time T_a , and a certain processor request rate: how do we model the bandwidth available from these memory modules, and how do we compute the overall effective access time? Clearly, the modules in low-order interleave are the only ones that can contribute to the bandwidth, and hence they determine m . From the memory system's point of view, it really does not matter whether the processor system consists of n processors, each making one request every memory cycle (i.e., one per T_c), or one processor with n requests per T_c , so long as the statistical distribution of the requests remains the same. Thus, to a first approximation, the analysis of the memory system is equally applicable to the multiprocessor system or the superscalar processor. The request rate, defined as n requests per T_c , is called the offered request rate, and it represents the peak demand that the non-cached processor system has on the main memory system.




Figure 4.28 Finding simple processor equivalence.

4.16.1 Models of Multiple Simple Processors and Memory

In order to develop a useful memory model we need a model of the processor. For our analysis we model a single processor as an ensemble of multiple simple processors. Each simple processor issues a request as soon as its previous request has been satisfied. Under this model, we can vary the number of processors and the number of memory modules and maintain the address request/data supply equilibrium. To convert the single processor model into an equivalent multiple processor, the designer must determine the number of requests to the memory module per module service time, $T_s = T_c$.

A simple processor makes a single request and waits for a response from memory. A *pipelined processor* makes multiple requests for various buffers before waiting for a memory response. There is an approximate equivalence between n simple processors, each requesting once every T_s , and one pipelined processor making n requests every T_s (Figure 4.28).

In the following discussion, we use two symbols to represent the bandwidth available from the memory system (the achieved bandwidth):

1. B : the number of requests that are serviced each T_s . Occasionally, we also specify the arguments that B takes on, e.g. $B(m, n)$ or $B(m)$.
2. Bw : the number of requests that are serviced per second: $Bw = B/T_s$.

To translate this into cache based systems, the service time, T_s , is the time that the memory system is busy managing a cache miss. The number of memory modules, m , is the maximum number of cache misses that the memory system can handle at one time and n is the total number of request per T_s . This is the total number of expected misses per processor per T_s multiplied by the number of processors making requests.

4.16.2 The Strecker - Ravi Model

This is a simple yet useful model for estimating contention. The original model was developed by Strecker [20] and independently by Ravi [16]. It assumes that there are n simple processor requests made per memory cycle and there are m memory modules. Further, we assume that there is no bus contention. The Strecker model assumes that

the memory request pattern for the processors is uniform and the probability of any one request to a particular memory module is simply $1/m$. The key modeling assumption is that the state of the memory system at the beginning of the cycle is not dependent upon any previous action on the part of the memory—hence, not dependent upon contention in the past (i.e., Markovian). Unserved requests are discarded at the end of the memory cycle.

The following modeling approximations are made:

1. A processor issues a request as soon as its previous request has been satisfied.
2. The memory request pattern from each processor is assumed to be uniformly distributed; i.e., the probability of any one request being made to a particular memory module is $1/m$.
3. The state of the memory system at the beginning of each memory cycle (i.e., which processors are awaiting service at which modules) is ignored by assuming that all unserviced requests are discarded at the end of each memory cycle and that the processors randomly issue new requests.

Analysis:

Let the average number of memory requests serviced per memory cycle be represented by $B(m, n)$. This is also equal to the average number of memory modules busy during each memory cycle. Looking at events from any given module's point of view during each memory cycle, we have:

$$\text{Prob (a given processor does not reference the module)} = (1 - 1/m)$$

$$\begin{aligned} \text{Prob (no processor references the module)} &= \text{Prob (the module is idle)} \\ &= (1 - 1/m)^n \\ \text{Prob (the module is busy)} &= 1 - (1 - 1/m)^n \\ B(m, n) = \text{average number of busy modules} &= m(1 - (1 - 1/m)^n) \end{aligned}$$

The achieved memory bandwidth is less than the theoretical maximum due to contention. By neglecting congestion in previous cycles, this analysis results in an optimistic value for the bandwidth. Still it is a simple estimate that should be used conservatively.

It has been shown by Bhandarkar [12] that $B(m, n)$ is almost perfectly symmetrical in m and n . He exploited this fact to develop a more accurate expression for $B(m, n)$, which is

$$B(m, n) = K \left[1 - (1 - 1/K)^l \right],$$

where $K = \max(m, n)$ and $l = \min(m, n)$.

We can use this to model a typical processor ensemble.

EXAMPLE 4.2

Suppose we have a two processor die system sharing a common memory. Each processor die is dual core with the two processors (four processors total) sharing a 4 MB level 2 cache. Each processor makes 3 memory references per cycle and the clock rate is 4 GHz. The L2 cache has a miss rate of 0.001 misses per reference. The memory system has an average T_S of 24 ns including bus delay.

We can ignore the details of the level 1 caches by inclusion. So each processor die creates 6×0.001 memory references per cycle or 0.012 references for both cycles. Since there are 4×24 cycles in a T_S , we have $n = 1.152$ processor requests per T_S . If we design the memory system to manage $m = 4$ requests per T_S , we compute the performance as

$$B(m, n) = B(4, 1.152) = 0.81.$$

The relative performance is

$$P_{\text{rel}} = \frac{B}{n} = \frac{0.81}{1.152} = 0.7.$$

Thus the processor can only achieve 70% of its potential due to the memory system. To do better we need either a larger level 2 cache (or a level 3 cache) or a much more elaborate memory system ($m = 8$).



4.16.3 Interleaved Caches

Interleaved caches can be handled in a manner analogous to interleaved memory.

EXAMPLE 4.3

An early Intel Pentium™ processor had an 8-way interleaved data cache. It makes two references per processor cycle. The cache has the same cycle time as the processor. For the Intel instruction set,

$$\text{Prob (data references per instruction)} = 0.6.$$

Since the Pentium tries to execute two instructions each cycle, we have

$$n = 1.2,$$


$$m = 8.$$

Using Strecker's model, we get:

$$B(m, n) = B(8, 1.2) = 1.18.$$

The relative performance is

$$P_{\text{rel}} = \frac{B}{n} = \frac{1.18}{1.2} = 0.98,$$



i.e., the processor slows down by about 2% due to contention.



4.17 Conclusions

Cache provides the processor with a memory access time significantly faster than the memory access time. As such the cache is an important constituent in the modern processor. The cache miss rate is largely determined by the size of the cache but any estimate of miss rate must consider the cache organization, the operating system, the system's environment, and I/O effects. As cache access time is limited by size, multi level caches are a common feature of on die processor designs.

On-die memory design seems to be relatively manageable especially with the advent of embedded DRAM but off die memory design is an especially difficult problem. The primary objective of such designs is capacity (or size); but large memory capacity and pin limitations necessarily imply slow access times. Even if die access is fast, the system's overhead, including bus signal transmission, error checking, and address distribution, add significant delay. Indeed, these overhead delays have increased relative to decreasing machine cycle times. Faced with hundred cycle memory access time, the designer can provide adequate memory bandwidth to match the request rate of the processor only by very large multi level cache.

4.18 Problem Set

1. A 128^{KB} cache has 64^{B} lines, 8^{B} physical word, 4^{KB} pages, and is four-way set associative. It uses copyback (allocate on write) and LRU replacement. The processor creates 30-bit (byte-addressed) virtual addresses that are translated into 24-bit (byte-addressed) real byte addresses (labeled A_0 – A_{23} , from least to most significant).
 - (a) Which address bits are unaffected by translation ($V = R$)?
 - (b) Which address bits are used to address the cache directories?
 - (c) Which address bits are compared to entries in the cache directory?
 - (d) Which address bits are appended to address bits in (b) to address the cache array?
2. Show a layout of the cache in problem 1. Present detail as in Figures 4.5 through 4.7.
3. Plot traffic (in bytes) as a function of line size for a DTMR cache (CBWA, LRU) for:
 - (a) 4^{KB} cache.
 - (b) 32^{KB} cache.
 - (c) 256^{KB} cache.
4. Suppose we define the miss rate at which a copyback cache (CBWA) and a write-through cache (WTNWA) have equal traffic as the crossover point.
 - (a) For the DTMR cache, find the crossover point (miss rate) for 16^{B} , 32^{B} , and 64^{B} lines. To what cache sizes do these correspond?
 - (b) Plot line size against cache size for crossover.
5. The cache in problem 1 is now used with a 16^b line in a transaction environment ($Q = 20,000$).
 - (a) Compute the effective miss rate.
 - (b) Approximately what is the optimal cache size (the smallest cache size that produces the lowest achievable miss rate)?
6. In a two-level cache system, we have:
 - L1 size 8KB with 4w set assoc., 16B lines, and write through (no allocate on writes)
 - and
 - L2 size 64KB direct mapping, 64B lines, and copy back (with allocate on writes)

Suppose the miss in L1, hit in L2 delay is 3 cycles and the miss in L1, miss in L2 delay is 10 cycles. The processor makes 1.5 refr/I.

- (a) What are the L1 and L2 miss rates?
- (b) What is the expected CPI loss due to cache misses?

- (c) Will *all* lines in L1 always reside in L2? Why?
7. A certain processor has a two-level cache. L1 is 4KB direct-mapped, WTNWA. The L2 is 8KB direct-mapped, CBWA. Both have 16-byte lines with LRU replacement.
- Is it always true that L2 includes all lines at L1?
 - If the L2 is now 8KB 4-way set associative (CBWA), does L2 include all lines at L1?
 - If L1 is 4-way set associative (CBWA) and L2 is direct-mapped, does L2 include all lines of L1?
8. Suppose we have the following parameters for an L1 cache with 4KB and an L2 cache with 64KB.

The cache miss rate is:


4KB	0.10 misses per refr
64KB	0.02 misses per refr
1 refr/I	
3 cycles	L1 miss, L2 hit
10 cycles	total time L1 miss, L2 miss

What is the excess CPI due to cache misses?

9. A certain processor produces a 32-bit virtual address. Its address space is segmented (each segment is 1 megabyte maximum) and paged (512-byte pages). The physical word transferred to/from cache is 4 bytes.
- A TLB is to be used, organized set associative, 128×2 . If the address bits are labeled V_0-V_{31} for virtual address and R_0-R_{31} for real address, least to most significant:
- Which bits are unaffected by translation (i.e., $V_i = R_i$)?
 - If the TLB is addressed by the low-order bits of the portion of the address to be translated (i.e., no hashing), which bits are used to address the TLB?
 - Which virtual bits are compared to virtual entries in the TLB to determine whether a TLB hit has occurred?
 - As a minimum, which real address bits does the TLB provide?
10. For an 16KB integrated level 1 cache (direct mapped, 16B lines) and a 128KB integrated level 2 cache (2W set associative, 16B lines) find the solo and local miss rate for the level 2 cache.
11. A certain chip has area sufficient for an 16KB I-cache and a 16KB D-cache, both direct mapped. The processor has a virtual address of 32b, real address of 26b, and uses 4KB pages. It makes 1.0 I-refr/I and 0.5 D-refr/I. The cache miss delay is 10 cycles plus 1 cycle for each 4B word transferred in a line. The processor is stalled until the entire line is brought into the cache. The D-cache is CBWA, use dirty line ratio $w = 0.5$. For both caches the line size is 64B. Find:



- (a) The CPI lost due to I-misses and the CPI lost due to D-misses.
 - (b) For the 64B line, find the number of I- and D-directory bits and corresponding **rbe** (area) for both directories.
12. Find two recent examples of DDR3 devices and for these devices update the entries of Table 4.9 and 4.10.
 13. List all the operations that must be performed after a “not in TLB” signal. How would a designer minimize the “not in TLB” penalty?
 14. In Study 4.2 suppose we need relative performance of 0.8. Would this be achieved by interleaving at $m=8$?
 15. Update the timing parameters for the NAND based Flash memory described in Table 4.3.
 16. Compare recent commercially available Flash (NAND and NOR) with eDRAM.



Bibliography

- [1] G. McFarland, *CMOS Technology Scaling and its Impact on Cache Delay*, Ph.D. Thesis, Stanford University, June 1997.
- [2] ARM1136J(F)-S Processor Specifications, ARM11 Family, 2010.
- [3] Freescale e600 Core Product Brief, Rev.0, 2004.
- [4] Freescale MPC8544E PowerQUICC III Integrated Processor, Hardware Specifications, Rev.2, 2009.
- [5] Intel PXA27x overview, 2010.
- [6] Philips Nexpria PNX1700 connected media processor, 2007.
- [7] NXP LH7A404, 32-Bit System-on-Chip, Preliminary data sheet, July 2007.
- [8] Intel IOP333 I/O Processor Chipset, July 2005.
- [9] Digi, NetSilicon NS9775 Datasheet, Rev. C January 2005.
- [10] A. Agarwal, *Analysis of Cache Performance of Operating Systems and Multiprogramming*. PhD thesis, Computer Systems Laboratory, Stanford University, May 1987. Published as CSL-TR-87-332.
- [11] J.E. Barth, J.H. Dreibelbis, E.A. Nelson, D.L. Anand, G. Pomichter, P. Jakobsen, M.R. Nelms, J. Leach and G.M. Belansek, Embedded DRAM design and architecture for the IBM 0.11-m ASIC offering, *IBM J. Res. Dev.*, 46(6), 2002, pp. 675–689.
- [12] D.P. Bhandarkar, Analysis of memory interference in multiprocessors. *IEEE Transactions on Computers*, C-24(9):897–908, September 1975.
- [13] Embedded DRAM Comparison Charts, *IBM Microelectronics Presentation*, December 2003.
- [14] I.J. Haikala, Program behavior in memory hierarchies. PhD thesis (Technical Report A-1986-2), CS Dept., University of Helsinki, 1986.
- [15] S. Przybylski, M. Horowitz and J. Hennessy, Characteristics of performance-optimal multi-level cache hierarchies. In *Proceedings of the 16th Annual Symposium on Computer Architecture*, pp. 114–121, June 1989.
- [16] C.V. Ravi, On the bandwidth and interference in interleaved memory systems. *IEEE Transactions on Computers*, C-21:899–901, August 1972.
- [17] A.J. Smith, Cache memories. *Computing Surveys*, 14(3):473–530, September 1982.
- [18] A.J. Smith, Cache evaluation and the impact of workload choice. In *Proceedings of the 12th International Symposium on Computer Architecture*, pages 64–73, June 1985.
- [19] E.J. Prinz, G.L. Chindalore, K. Harber, C.M. Hong, C.B. Li and C.T. Swift, Sonos: An Embedded 90nm SONOS Flash EEPROM Utilizing Hot Electron Injection Programming and 2-Sided Hot Hole Injection Erase. *IEDM Conference record*, 2002.

- [20] W.D. Strecker, *Analysis of the Instruction Execution Rate in Certain Computer Systems.* PhD thesis, Carnegie-Mellon University, Pittsburgh, PA, 1970.
- [21] D. Fisch, A. Singh and G. Popov, Z-RAM Ultra-dense Memory for 90nm and Below. Hot Chips 18, August 2006.